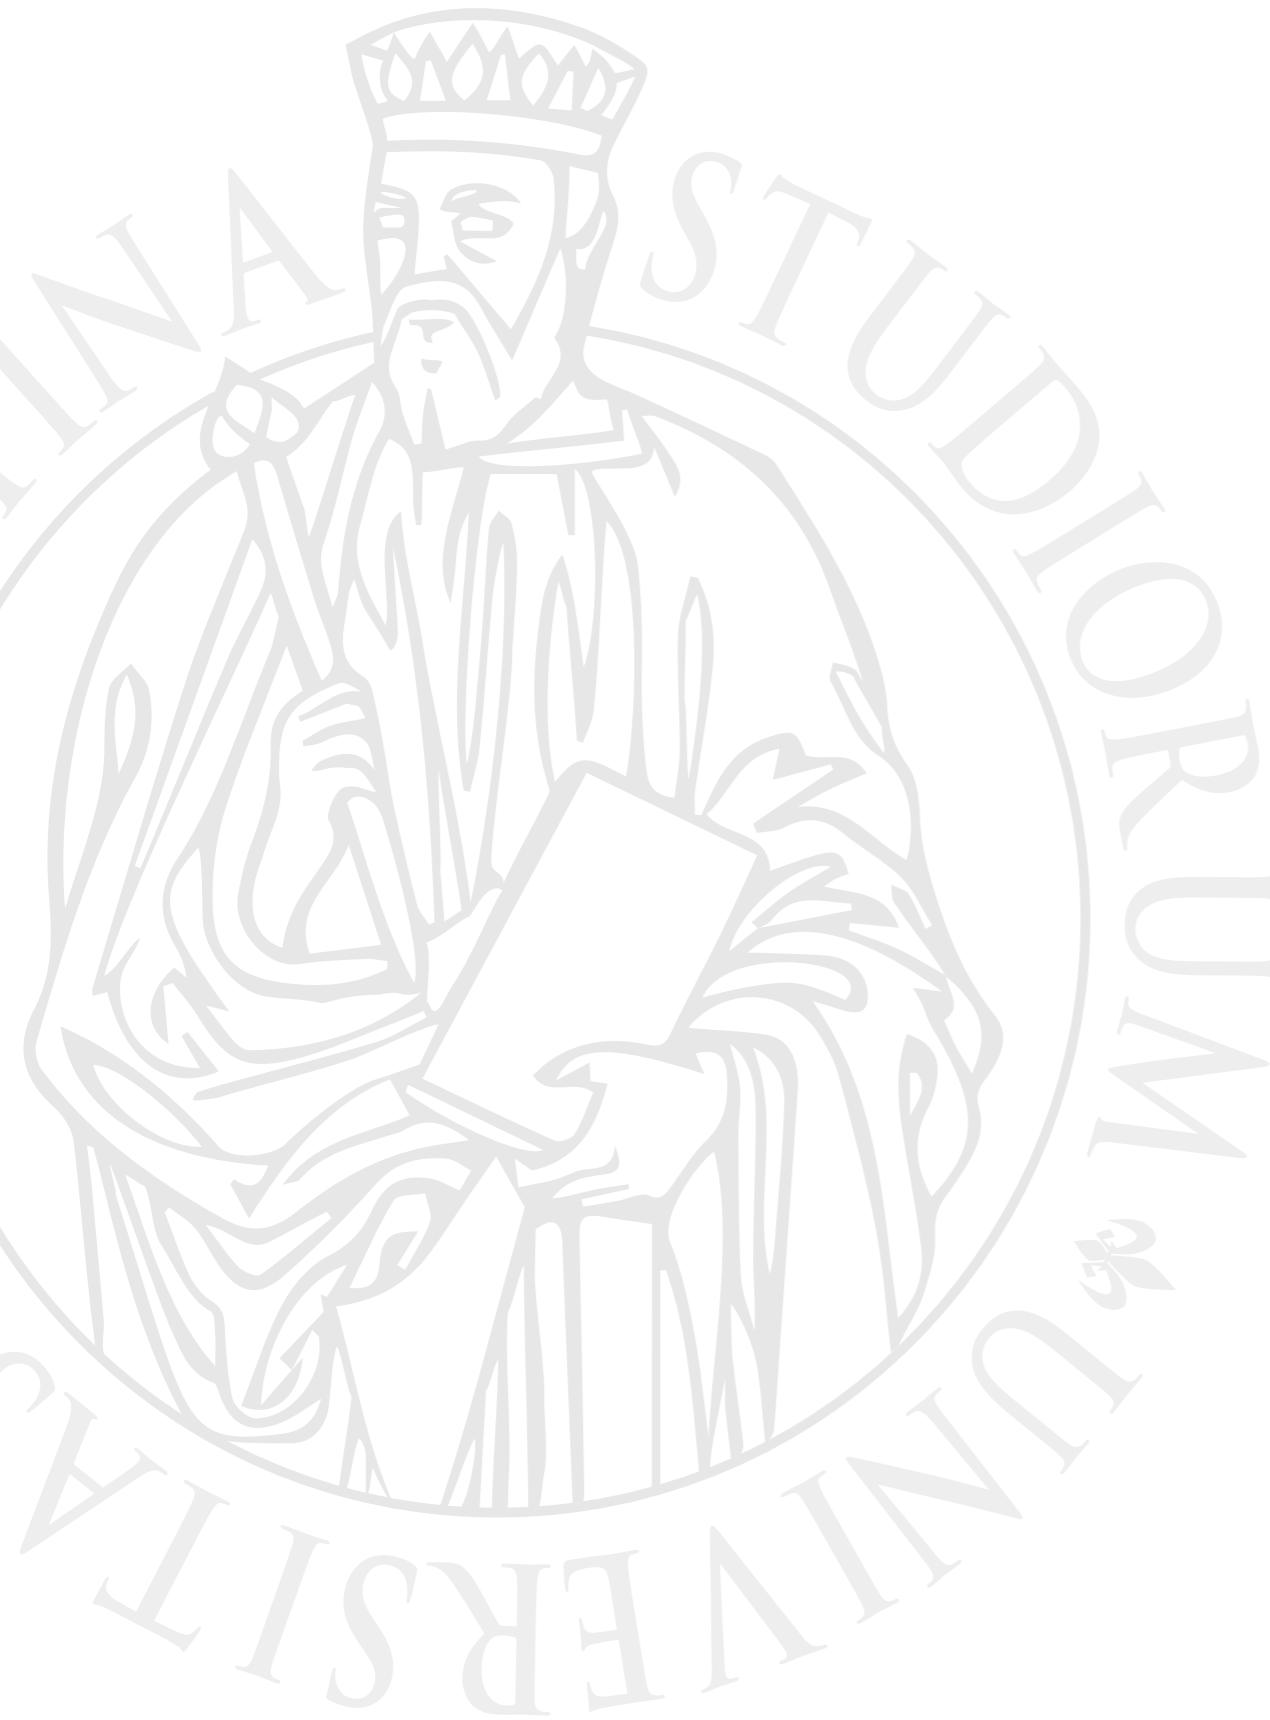




UNIVERSITÀ  
DEGLI STUDI  
FIRENZE



# **Mid-term programming assignments**

## **Parallel Programming**



# General guidelines

- The goal is to understand how to parallelize a program and evaluate the performance of the parallelization
- Assignments can be implemented individually or in pairs (only for 9-credit course)
  - Both team members must contribute equally and understand all aspects of the implementation

# Implementation requirements

- Implement both sequential and parallel versions
  - using OpenMP + vectorization techniques
- Use C/C++ as the primary language
- Focus on data parallelism, loop optimization, and proper synchronization
- Apply vectorization where applicable (consider SIMD operations)

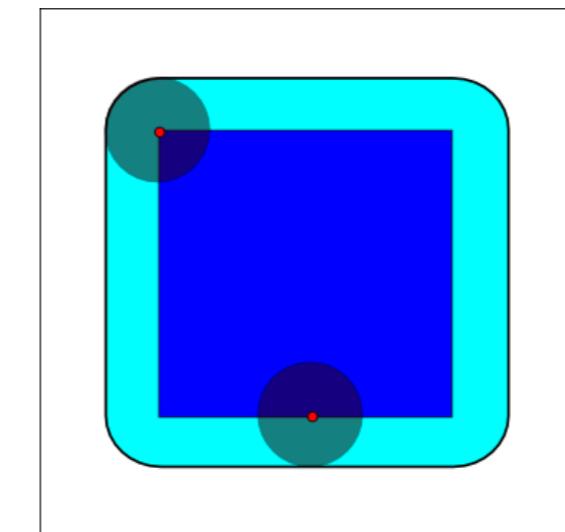
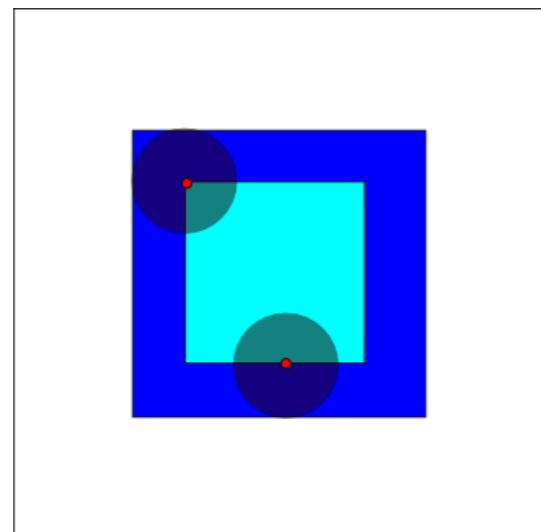
# Deliverables

- Source Code: Well-documented, clean code on public repository (GitHub/GitLab/Bitbucket)
- Technical Report: Comprehensive performance analysis (see detailed guidelines below)
- Presentation: Clear explanation of approach, results, and insights



# Morphological image processing

- Implement morphological operations (erosion, dilation, opening, closing) on grayscale or binary images.
  - Implement at least 2-3 morphological operations
  - Test on various image sizes (e.g.  $512 \times 512$ ,  $1024 \times 1024$ ,  $2048 \times 2048$ ,  $4096 \times 4096$ )
  - Consider different structuring element sizes
  - Apply on real images (medical imaging, satellite imagery, etc.)



# Morphological image processing

## References

- [https://en.wikipedia.org/wiki/Mathematical\\_morphology](https://en.wikipedia.org/wiki/Mathematical_morphology) and [https://homepages.inf.ed.ac.uk/rbf/CVonline/LOCAL\\_COPIES/GASTERATOS/SOFT/2.htm](https://homepages.inf.ed.ac.uk/rbf/CVonline/LOCAL_COPIES/GASTERATOS/SOFT/2.htm)
- Example applications: [https://docs.opencv.org/4.x/d9/d61/tutorial\\_py\\_morphological\\_ops.html](https://docs.opencv.org/4.x/d9/d61/tutorial_py_morphological_ops.html)

## Datasets

- Medical images: <https://www.kaggle.com/datasets> (search "medical imaging", "X-ray", "MRI")
- Berkeley Segmentation Dataset: <https://www2.eecs.berkeley.edu/Research/Projects/CS/vision/bsds/>
- DIV2K (high-resolution images): <https://data.vision.ee.ethz.ch/cvl/DIV2K/>

# Morphological image processing

Parallelization Opportunities:

- Pixel-level parallelism for operations
- Sliding window parallelism
- Pipeline parallelism for sequential operations



# Bigrams / trigrams

- Compute histograms of bigrams/trigrams on texts (eg. Wikipedia / Gutenberg project documents)
  - Consider bigrams/trigrams of characters and words
    - Handle tokenization and normalization (lowercase, punctuation)
  - Generate frequency distributions and statistics
    - If needed re-use more and more times the collected data to create a large enough corpus

# Bigrams / trigrams

Considerations:

- Token splitting strategies
- Memory-efficient data structures (hash maps, tries)

Links / datasets

- N-gram tutorial: <https://web.stanford.edu/~jurafsky/slp3/3.pdf>
- Project Gutenberg: <https://www.gutenberg.org/>
- Wikipedia dumps: <https://dumps.wikimedia.org/>
- Common Crawl: <https://commoncrawl.org/>
- Large text corpus: <https://wortschatz.uni-leipzig.de/en/download>



# Bigrams / trigrams

Parallelization Opportunities:

- Document-level parallelism
- Chunk-based text processing
- Parallel hash table updates with atomic operations or locks



# ANN retrieval

- Goal: find k-nearest neighbors (k-NN overview: <https://scikit-learn.org/stable/modules/neighbors.html>) for query vectors in high-dimensional space.
  - Consider high dimensionality vectors (eg.  $\geq 128$ )
  - Support various distance metrics (Euclidean, cosine similarity, Manhattan)
  - Test with k values: 1, 10, 100
- Given a specified query find the *k nearest neighbors* in a database of vectors
  - Can use LSH for approximate indices, eg. using mlpack or LSHkit
  - Can get sample datasets from <http://corpus-texmex.irisa.fr>

# ANN retrieval

Parallelization Opportunities:

- Parallel/vectorized distance computations
- Parallel search across database partitions
- Parallel hash table construction for LSH



# Password decryption

- Decrypt passwords encrypted using crypt (DES)
  - Consider 8 characters lengths in the set [a-zA-Z0-9./].  
Can focus on specific patterns (dates, common passwords, dictionary words)
  - Library funds may be deprecated or not available. E.g. CUDA does not provide the crypt C-library function. Must use a special implementation (check Moodle).
  - It's a (slowed down) search problem.
  - Can choose to attack a specific password in a list (in different positions), or decrypt a full list of passwords (in this case reduce the search space)
  - Assume salt is known.

# Password decryption

Datasets:

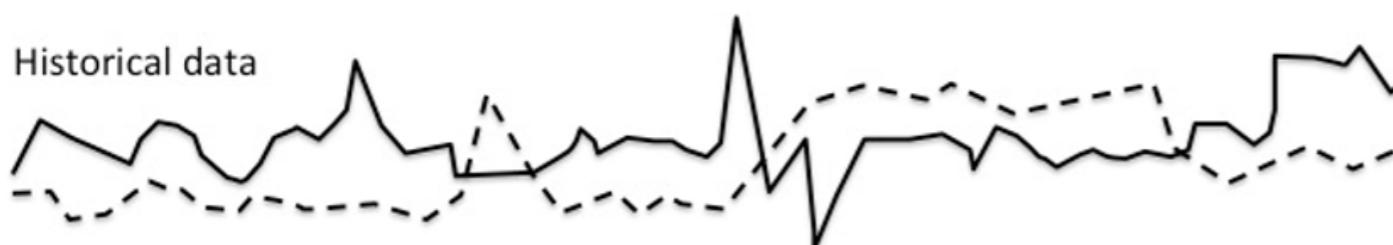
- Use common password lists (e.g., RockYou dataset samples)
- Generate encrypted passwords for testing
- Focus on realistic password patterns to reduce search space

Parallelization Opportunities:

- Search space partitioning
- Task parallelism for multiple passwords
- Load balancing strategies for irregular workloads

# Pattern recognition in time series

- Search a given time series within a larger and longer set of time series
  - Implement sliding window search
  - Use SAD (Sum of Absolute Differences) or similar metrics (DTW optional)



$\alpha_1$	$\alpha_2$	$\alpha_3$	$\alpha_4$	$\alpha_5$	$\alpha_6$	$\alpha_7$	$\alpha_8$	$\alpha_9$	$\alpha_{10}$	$\alpha_{11}$	$\alpha_{12}$	$\alpha_{13}$	$\alpha_{14}$	$\alpha_{15}$	$\alpha_{16}$
$\beta_1$	$\beta_2$	$\beta_3$	$\beta_4$	$\beta_5$	$\beta_6$	$\beta_7$	$\beta_8$	$\beta_9$	$\beta_{10}$	$\beta_{11}$	$\beta_{12}$	$\beta_{13}$	$\beta_{14}$	$\beta_{15}$	$\beta_{16}$



$\gamma_1$	$\gamma_2$	$\gamma_3$	$\gamma_4$	$\gamma_5$	$\gamma_6$
$\delta_1$	$\delta_2$	$\delta_3$	$\delta_4$	$\delta_5$	$\delta_6$

$$|\alpha_1 - \gamma_1| + \dots + |\alpha_6 - \gamma_6| + \\ |\beta_1 - \delta_1| + \dots + |\beta_6 - \delta_6|$$

$$|\alpha_{11} - \gamma_1| + \dots + |\alpha_{16} - \gamma_6| + \\ |\beta_{11} - \delta_1| + \dots + |\beta_{16} - \delta_6|$$

# Pattern recognition in time series

- Handle multiple pattern queries
- Work with time series of length 1.000-100.000 points

Datasets:

- UCR Time Series Archive
- Stock market data
- Sensor data (IoT, weather)
- Generate synthetic data: [https://github.com/cyrilou242/  
mockseries](https://github.com/cyrilou242/mockseries)

# Pattern recognition in time series

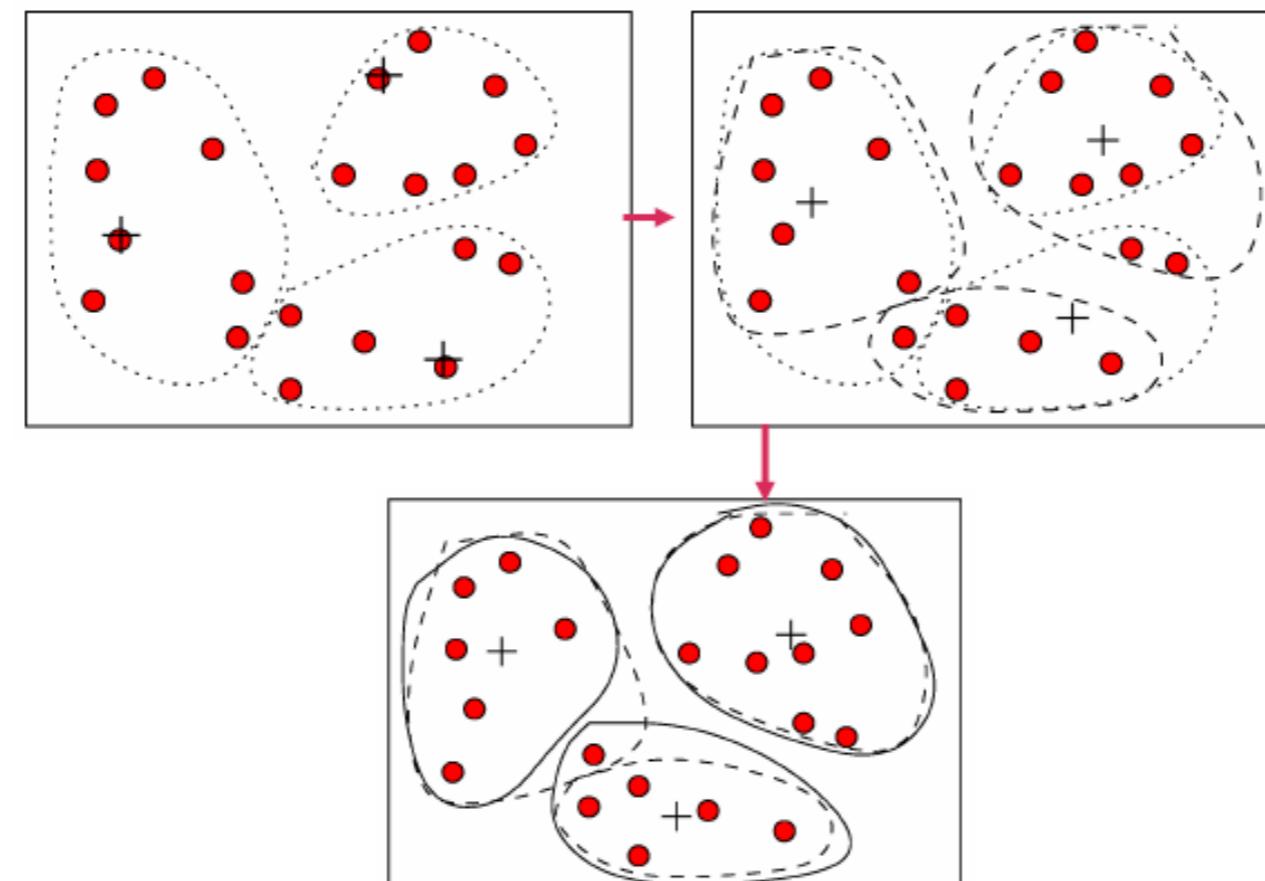
Parallelization Opportunities:

- Parallel sliding window evaluation
- Multiple pattern search in parallel
- Data parallelism across time series database



# k-means

- Implement a parallel version of k-means clustering
  - There's no need to implement K-means++ centroid assignments
  - Careful with experiment reproducibility: use fixed iterations or the same starting points
  - Feel free to chose to operate on 2D or larger dimensionality vectors





# k-means

Datasets:

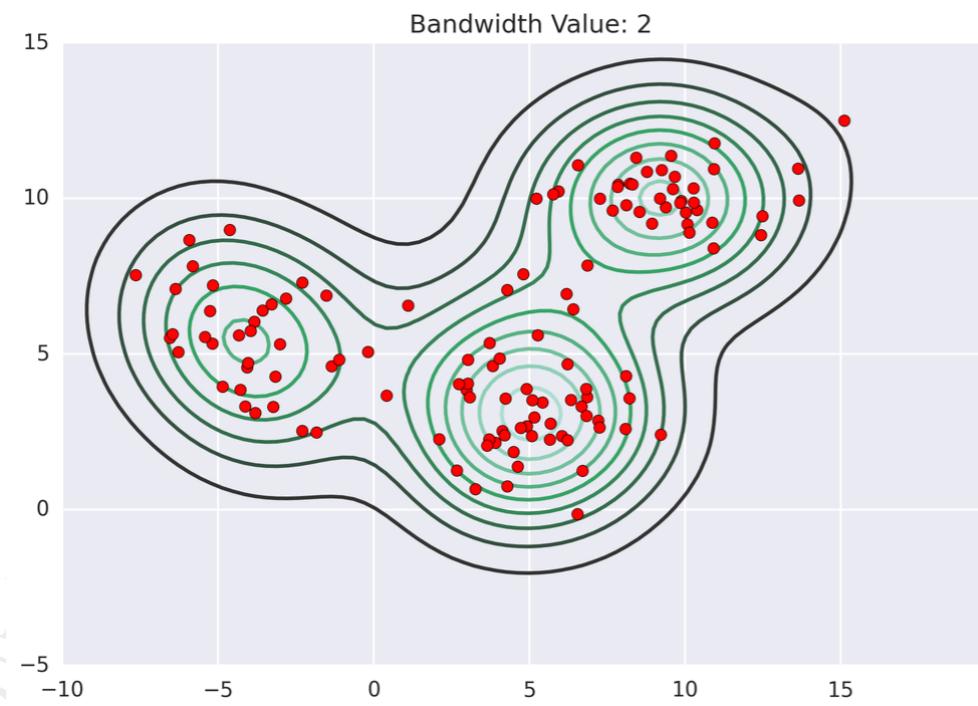
- Iris, Wine, Digits datasets
- Image segmentation (pixel clustering)
- Synthetic blob datasets, e.g. <https://python-course.eu/machine-learning/artificial-datasets-with-scikit-learn.php>

Parallelization Opportunities:

- Parallel distance computation for assignment step
- Parallel centroid update
- Reduction operations for cluster statistics
- AoS vs SoA

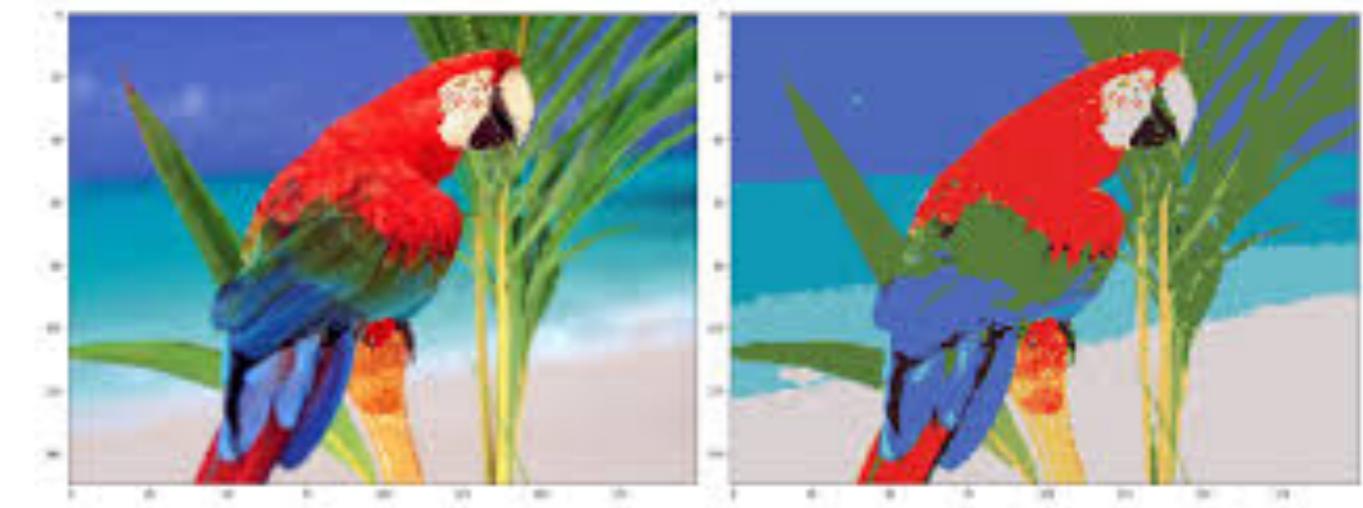
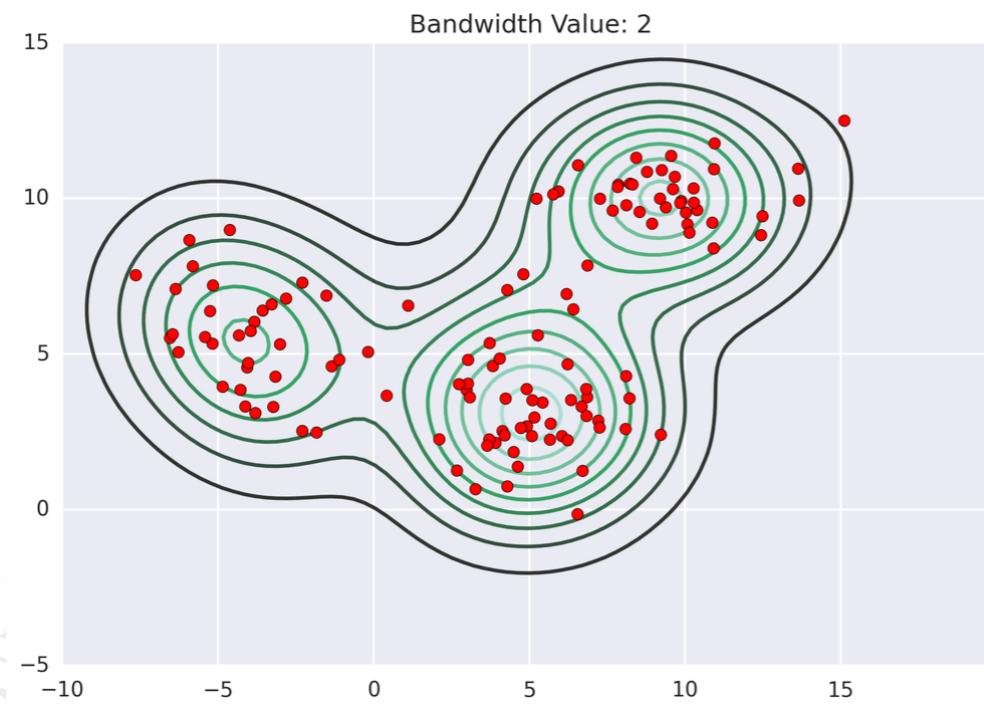
# Mean shift clustering

- Implement a parallel version of mean-shift clustering.  
This clustering doesn't require to specify the number of desired clusters (it uses KDE, so needs a bandwidth parameter) but its complexity is  $O(N^2)$ 
  - Test various bandwidth values
  - Can be used to segment images



# Mean shift clustering

- Implement a parallel version of mean-shift clustering.  
This clustering doesn't require to specify the number of desired clusters (it uses KDE, so needs a bandwidth parameter) but its complexity is  $O(N^2)$ 
  - Test various bandwidth values
  - Can be used to segment images



# Mean shift clustering

Datasets:

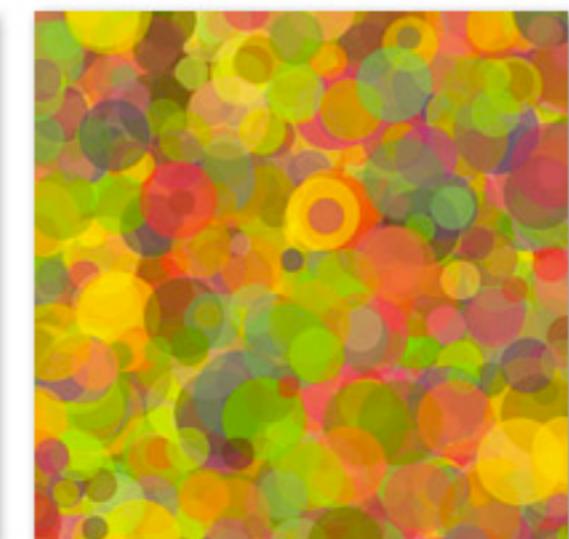
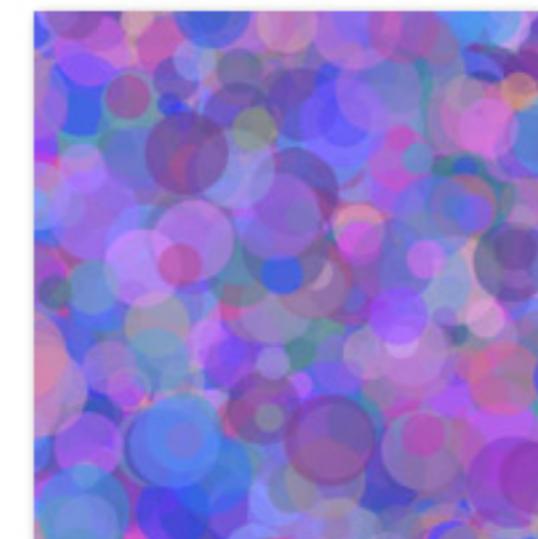
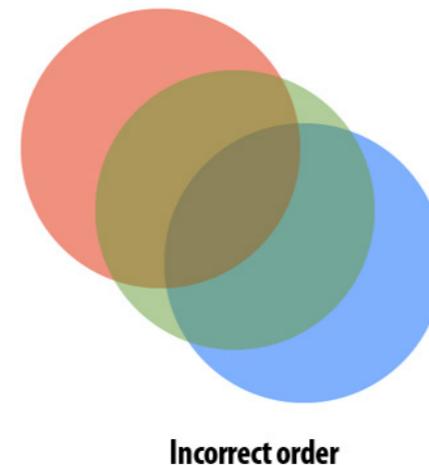
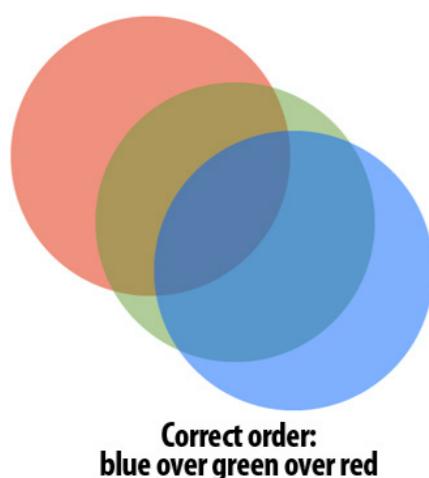
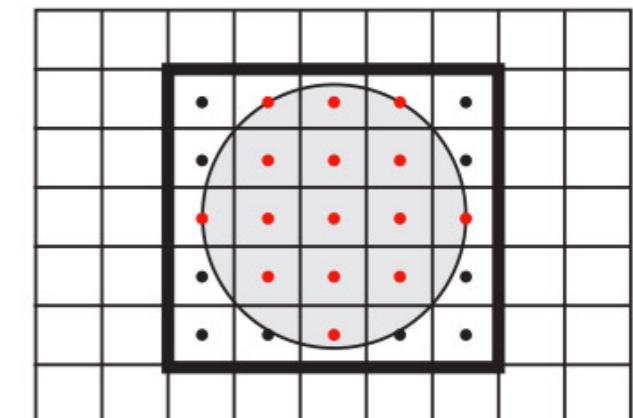
- Berkeley Segmentation Dataset
- COCO dataset samples
- Synthetic 2D/3D point clouds

Parallelization Opportunities:

- Parallel mean computation for each point
- Parallel kernel density estimation
  - Mode detection parallelism
-

# Renderer

- Implement a renderer that draws (semi) transparent circles (or other shapes). Circles have 3D coordinates and the order along Z axis matters for the correct rendering
  - A pixel belongs to a circle if its center is within the circle.





# Renderer

- Handle transparency/alpha blending
- Sort by Z-axis for correct rendering
- A pixel belongs to a shape if its center is within the shape

Output:

- Generate images at various resolutions (e.g.  $512 \times 512$ ,  $1024 \times 1024$ ,  $2048 \times 2048$ )
- Support hundreds to thousands of shapes

Parallelization Opportunities:

- Pixel-level parallelism
- Shape-level parallelism with synchronization
- Tile-based rendering



# Image composition

- Goal: parallel alpha composition of multiple images (variation of Renderer assignment).
- Implement a parallel alpha blending, supporting multiple overlay images
- Useful to create image augmentation to train Convolutional Neural Networks, e.g. for detection (see Blend augmentation in ImgAug library)
- Can parallelize applying the same object to different images, or in different positions in the same image

# Image composition

Datasets:

- COCO dataset
- Pascal VOC
- Custom object cutouts with alpha channels

Parallelization Opportunities:

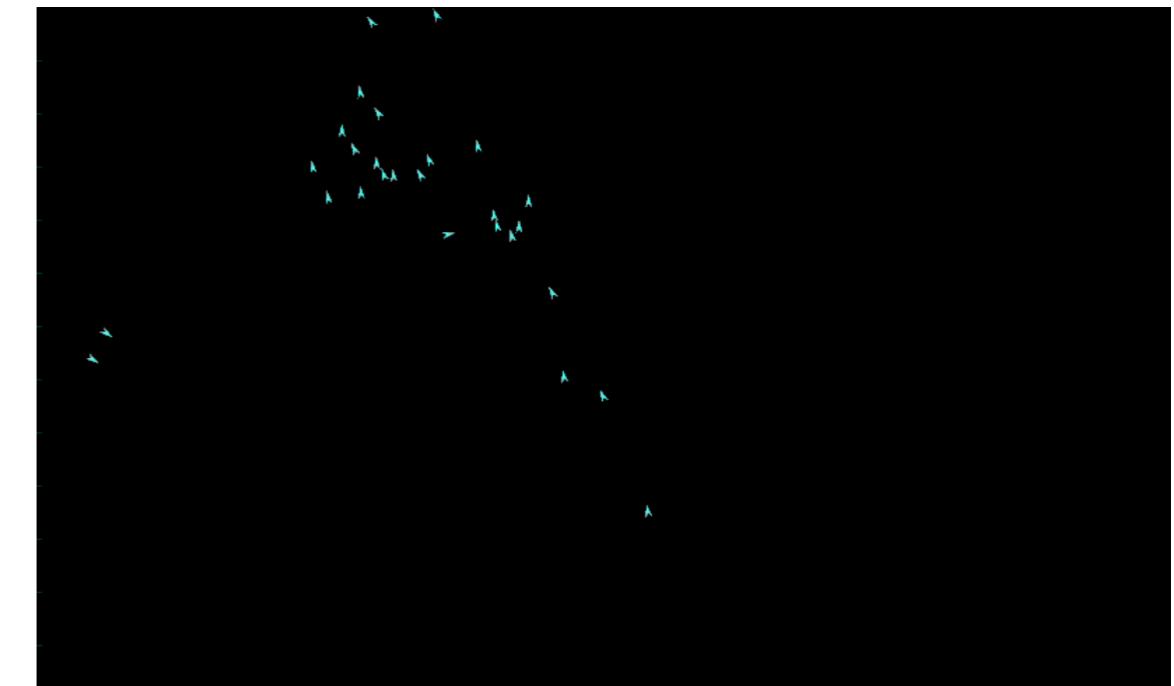
- Parallel composition of image pairs
- Batch processing of augmentations
- Pixel-level parallelism



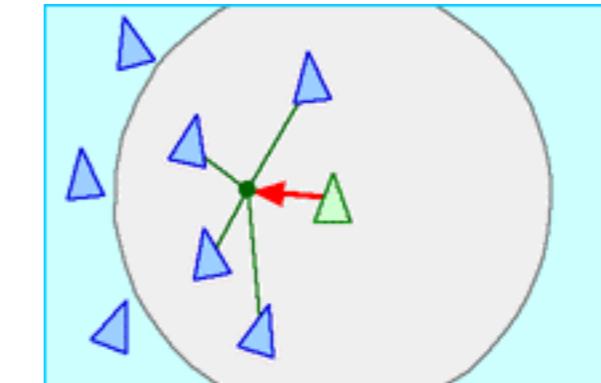
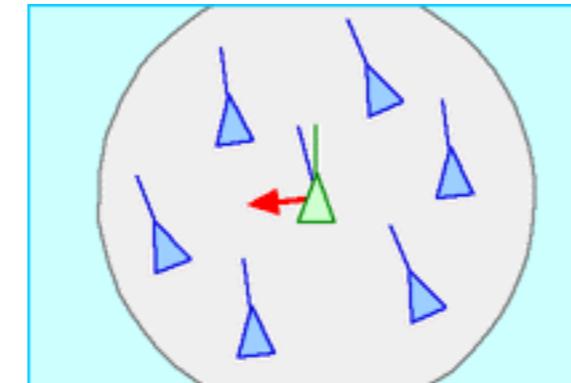
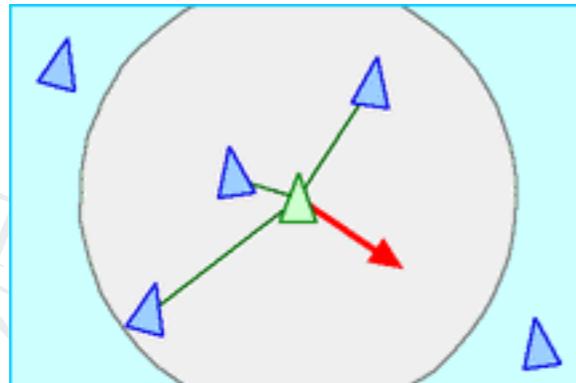
# Boids

- Simulate the behavior of a flock / crowd.

- Boids is an artificial life program which simulates the flocking behaviour of birds, and related group motion.



- Check [https://vanhunteradams.com/Pico/  
Animal\\_Movement/Boids-algorithm.html](https://vanhunteradams.com/Pico/Animal_Movement/Boids-algorithm.html)

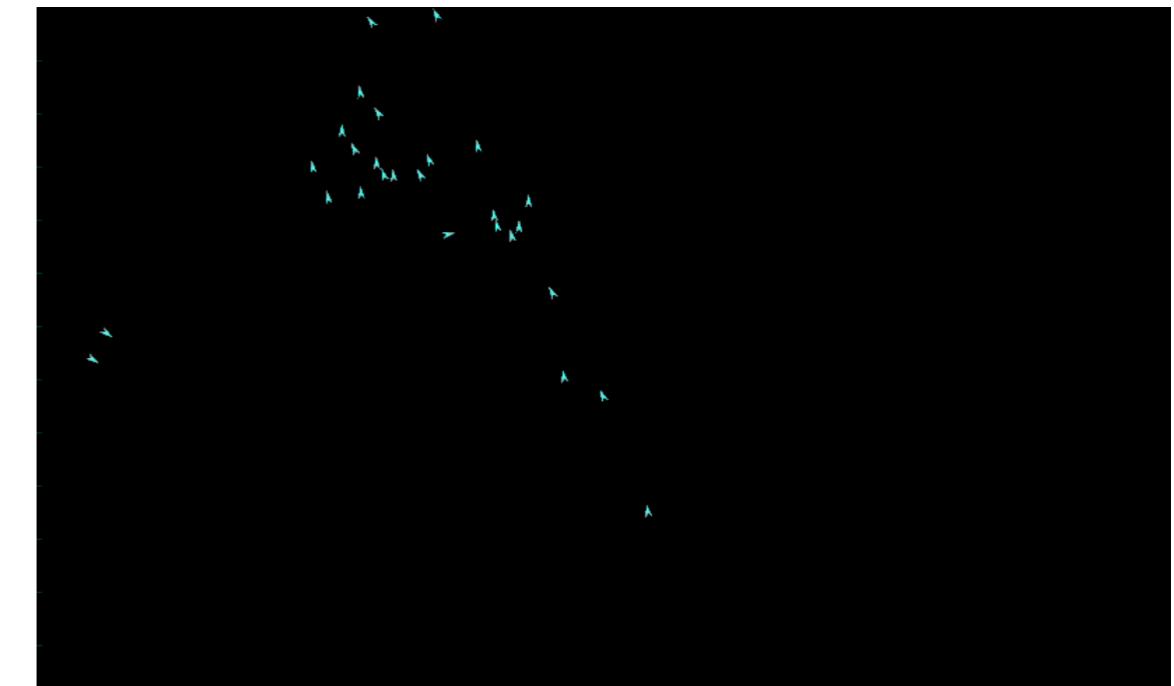




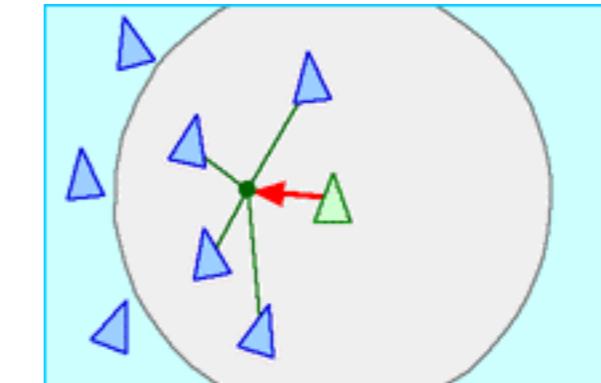
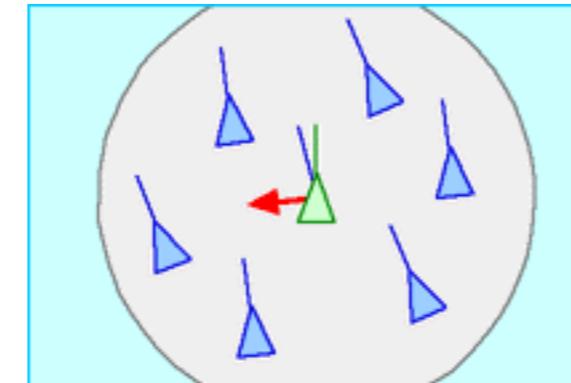
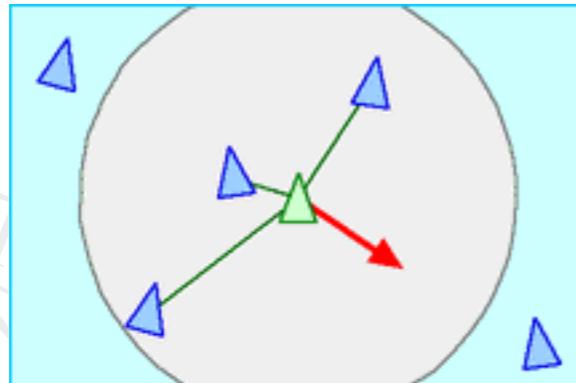
# Boids

- Simulate the behavior of a flock / crowd.

- Boids is an artificial life program which simulates the flocking behaviour of birds, and related group motion.



- Check [https://vanhunteradams.com/Pico/  
Animal\\_Movement/Boids-algorithm.html](https://vanhunteradams.com/Pico/Animal_Movement/Boids-algorithm.html)





# Boids

- Implement three basic rules: separation, alignment, cohesion
- Simulate 1000-10000+ agents
- Run for 1000+ time steps
- Optional: visualization of simulation

Parallelization Opportunities:

- Parallel neighbor finding (spatial partitioning)
- Parallel rule computation per agent
- Parallel position/velocity updates

# Non-Maximum Suppression (NMS) for Object Detection

- Goal: Implement parallel NMS to filter overlapping bounding box predictions in object detection pipelines.
- Applications:
  - Post-processing in object detection (YOLO, Faster R-CNN, SSD)
  - Any detection pipeline producing multiple overlapping predictions
- Critical component in many modern object detectors
  - Recent methods are NMS-free
- Tutorial: <https://learnopencv.com/non-maximum-suppression-theory-and-implementation-in-pytorch/>



# Non-Maximum Suppression (NMS) for Object Detection

- Challenges:
  - Sequential nature of greedy NMS (each iteration depends on previous)
  - Race conditions when multiple threads try to suppress same box
  - Load balancing (number of boxes per class/region varies)
  - Memory access patterns (scattered box data)
- Requirements:
  - Standard NMS (greedy sequential algorithm as baseline)
  - Implement parallel variants:
    - Parallel NMS with atomic operations
    - Soft-NMS (weighted score reduction instead of hard suppression)
    - Class-agnostic and class-specific NMS
- Process detection outputs with varying numbers of boxes:
  - 100, 1000, 10,000, 100,000 candidate boxes
  - Support multiple IoU (Intersection over Union) thresholds: 0.3, 0.5, 0.7
  - Handle multi-class scenarios (80+ classes like COCO)

# Non-Maximum Suppression (NMS) for Object Detection

- Parallel NMS Strategies:
  - Optimistic Parallel NMS:
    - Process boxes in parallel
    - Use atomic operations to mark suppressed boxes
    - May miss some suppressions but faster
  - Soft-NMS:
    - Instead of removing boxes, decay their scores
    - All boxes processed independently
    - More amenable to parallelization
  - Sorted List Approach:
    - Sort all boxes by confidence
    - Process in chunks, each chunk in parallel
    - Synchronize between chunks
- Spatial Partitioning:
  - Divide image into grid cells
  - Apply NMS independently in each cell
  - Handle boundary cases carefully



# Non-Maximum Suppression (NMS) for Object Detection

## Datasets:

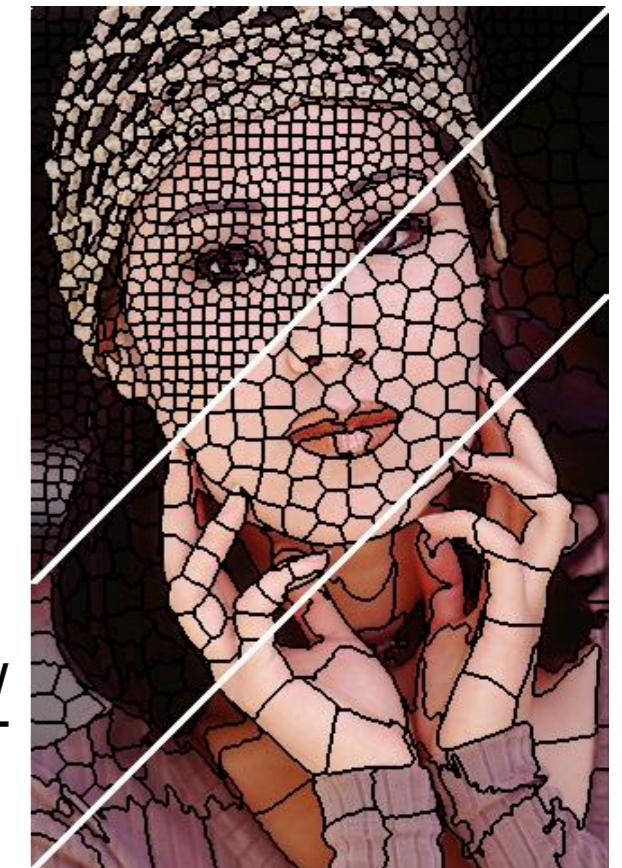
- COCO detection results (download pre-computed detections)
- Pascal VOC detection outputs
- Generate synthetic detection boxes with known ground truth
- Real detector outputs from pre-trained models (YOLO, Faster R-CNN)

## Good Practices:

- Validate correctness against sequential NMS (should get identical or very similar results)
- Measure precision/recall impact of parallel approximations
- Profile IoU computation (often the bottleneck)

# SLIC superpixel

- Implement the SLIC algorithm (k-means in CIE-Lab+XY) to segment images into K superpixels.
- SLIC (Simple Linear Iterative Clustering) reduces an image to a few hundred/thousand superpixels - i.e. compact, edge-aware regions useful in many CV pipelines. The algorithm is iterative, data-parallel, and memory-bandwidth bound
- SLIC superpixel tech report: <https://infoscience.epfl.ch/entities/publication/2dd26d47-3d00-43eb-9e31-4610db94a26e>
- Comparison of segmentation and superpixel algorithms: [https://scikit-image.org/docs/0.25.x/auto\\_examples/segmentation/plot\\_segmentations.html](https://scikit-image.org/docs/0.25.x/auto_examples/segmentation/plot_segmentations.html)



# SLIC superpixel

- Datasets:
  - Berkeley Segmentation (BSDS500 - <https://www.kaggle.com/datasets/balraj98/berkeley-segmentation-dataset-500-bsds500>),  
COCO samples, other high-res photos.

## Parallelization Opportunities:

- Parallel assignment step over pixels.
- Parallel cluster update using reductions.
- Experiment with tile/block processing to improve cache locality.
- SoA vs. AoS: how to apply it to an image ?



UNIVERSITÀ  
DEGLI STUDI  
FIRENZE



# Implementation guidelines



# Dataset

Finding Appropriate Datasets:

- Size matters: Ensure dataset is large enough to show meaningful parallelization benefits
- Rule of thumb: computation time should be  $\geq 10$  seconds sequential for good measurements
- If dataset is too small, replicate or synthesize larger versions

Realism: Use real-world datasets when possible for practical relevance

- UCI Machine Learning Repository
- Kaggle datasets
- OpenML datasets
- Government open data portals
- Synthetic data: When real data is unavailable, generate synthetic datasets

Data Preprocessing:

- Document all preprocessing steps
- Ensure preprocessing doesn't dominate computation time
- Consider whether preprocessing should be parallelized separately

# Parallelization Parameters

Number of Threads:

- Test: 1, 2, 4, 8, 16, 32 threads (up to 2× physical cores)
- Identify optimal thread count
- Observe superlinear speedup (if any) or speedup saturation

Scheduling Strategies (OpenMP):

- `schedule(static)`: Equal-sized chunks, predetermined
- `schedule(dynamic)`: Dynamic load balancing
- `schedule(guided)`: Decreasing chunk sizes
- `schedule(static, chunk_size)`: Experiment with chunk sizes
- Test which works best for your workload distribution

Chunk Sizes (when using static or dynamic scheduling):

- Too small: excessive overhead
- Too large: poor load balancing
- Test: automatic, manual specification (powers of 2, problem-dependent)

# Parallelization Parameters

Memory Access Patterns:

- Compare row-major vs column-major access
- Test effect of data structure layout (AoS vs SoA)
- EXTRA: Analyze cache hit rates

Synchronization Strategies:

- Critical sections vs atomic operations vs reduction clauses
- False sharing mitigation (padding, alignment)
- Lock granularity (fine-grained vs coarse-grained)

Vectorization Parameters:

- Test with/without vectorization flags
- EXTRA: Analyze vectorization reports



# Performance metrics

Execution Time:

- Wall-clock time (most important for users)
- CPU time (sum of all thread times)
- Measure multiple runs (minimum 5) and report statistics (mean, std, min, max)

Speedup:

- Speedup( $p$ ) =  $T_{\text{sequential}} / T_{\text{parallel}}(p)$
- Report speedup curves (speedup vs thread count)

Scalability

- Strong scaling: fixed problem size, increasing threads
- EXTRA: Weak scaling: problem size grows proportionally with threads



# Performance metrics

## Timing

- Use high-resolution timers: `omp_get_wtime()`,  
`std::chrono::high_resolution_clock`
- Warm-up runs: discard first 1-2 runs to eliminate cold-start effects
- Multiple measurements: run at least 5-10 times, report median or mean with confidence intervals
- Measure only computation, exclude I/O and initialization when possible

## System Configuration

- Document hardware: CPU model, core count, cache sizes, RAM, compiler version
- Close unnecessary background processes
- If possible disable CPU frequency scaling



# Technical report

## Abstract

- Brief problem description
- Approach summary
- Key results (best speedup achieved)

## Introduction

- Problem description and motivation
- Computational challenges
- Objectives and scope

## Algorithm Description

- Sequential algorithm explanation
- Identification of parallelizable sections

## Parallel Implementation

- Parallelization strategy
- OpenMP directives used (#pragma omp parallel for, reduction, etc.)
- Synchronization mechanisms
- Vectorization techniques applied
- Data structures and memory layout choices
- Challenges encountered and solutions

## Experimental Setup

- Hardware specifications
- Software environment (OS, compiler, flags)
- Datasets used (description, sizes, sources)
- Parameters tested
- Measurement methodology

# Technical report

## Results and Analysis

- Correctness validation: how you verified correctness
- Performance results:
- Speedup curves (speedup vs threads) with multiple input sizes
- Scalability analysis (strong/weak scaling)
- Effect of different parameters (scheduling, chunk size, etc.)
- Comparison of variants: different parallelization strategies
- Bottleneck analysis: what limits scalability?
- All plots should have clear labels, legends, and captions

## Discussion

- Interpretation of results
- Comparison with theoretical expectations (Amdahl's Law, Gustafson's Law)
- Overhead analysis
- Limitations and constraints
- Lessons learned
- Conclusions
- Summary of achievements
- Best configuration found
- Future work suggestions

## References

- Academic papers, documentation, dataset sources

# Presentation

- Duration: 10-12 min + 3-5 min. Questions/discussion

Structure:

- Problem Introduction (1-2 min)
  - What problem are you solving?
  - Why is it computationally challenging?
- Approach (3-4 min)
  - Sequential algorithm overview
  - Parallelization strategy
  - Key optimizations
- Results (4-5 min)
  - Speedup curves (main result)
  - Best configuration achieved

- Interesting findings or surprises
- Conclusions (1-2 min)
- Summary
- Lessons learned

Tips:

- Visualizations: Use clear graphs, avoid cluttered plots
- Focus: Don't explain every detail, highlight key insights. Probably you do not need to copy/paste code !
- Demo: If possible, show a quick live demo or image/video
- Rehearse: Practice timing
- Keep extra slides for discussion/answering details

# Code

- Version control: small, descriptive commits; tags for milestones.
- Documentation: short README, build/run instructions, parameter table, dataset citations.
- OpenMP official documentation:  
<https://www.openmp.org/specifications/>
- OpenMP tutorials:  
<https://computing.llnl.gov/tutorials/openMP/>



# Common pitfalls to avoid

- Measurement Errors:
  - Not running enough iterations for stable measurements
  - Including I/O time in performance measurements
  - Not accounting for cold cache effects
  - Comparing debug vs release builds
- Implementation Issues:
  - Race conditions due to missing synchronization
  - False sharing causing performance degradation
  - Incorrect use of private vs shared variables in OpenMP
  - Not validating correctness of parallel implementation
- Analysis Problems:
  - Testing only on small inputs that don't show parallelization benefits
  - Not exploring different thread counts systematically
  - Ignoring load imbalance issues
  - Not comparing different scheduling strategies
- Report Issues:
  - Missing hardware specifications
  - Plots without labels, legends, or captions
  - No discussion of why certain results occurred
  - Not explaining deviations from expected behavior

# Evaluation criteria

- Correctness (~25%)
  - Correct sequential implementation
  - Correct parallel implementation (validated against sequential)
  - Proper handling of edge cases
- Performance (~25%)
  - Achieved speedup
  - Efficiency of parallelization
  - Appropriate use of OpenMP features
  - Evidence of optimization effort
- Experimental Methodology (~20%)
  - Comprehensive parameter exploration
  - Proper measurement techniques
  - Statistical rigor
  - Variety of test cases
- Report/Presentation Quality (~20%)
  - Clear writing, organization and presentation
  - Thorough analysis and discussion
  - Quality of visualizations
  - Completeness
- Code Quality (~10%)
  - Readability and documentation
  - Repository organization